



Computer
Science

CSC380: Principles of Data Science

Course wrap-up 2

Xinchen Yu

Announcements

- Final exam
 - Time: Dec 13, 3:30 - 5:30pm
 - Location: C E Chavez Bldg, Rm 111 (same room)
 - What you can bring:
 - one letter size cheat sheet, you can use double sides
 - calculator (not necessary)
- Fill out SCS (<https://scsonline.oia.arizona.edu/>) – if 80% responses, will add 5 points to the homework with lowest grade (63% right now).

Announcements

- Grades on D2L are available by this Friday.
- Final project
 - Due this Friday by 11:59 pm
 - Evaluate on test set (1,086 instances)
- Final exam practice questions out on D2L->Content

Predictive Modeling and Classification

How to construct a decision tree


5

- Assign all training instances to the root of the tree. Set current node to root node.
- For each feature:
 - a. Partition all data instances at the node by the value of the feature.
 - b. Compute the accuracy from the partitioning.
- Identify feature that results in the highest accuracy. Set this feature to be the splitting criterion at the current node.


		Prereqs	Lecturer	HasLabs	
		↓	↓	↓	
Rating	Easy?	Alt?	Sys?	Thy?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y

Decision tree: accuracy


6




HasTakenPrereqs




SameLecturer



HasLabs



Suppose we place the node HasTakenPrereqs at the root. Set the prediction at each leaf node as the majority vote.



What is the train set accuracy now?

$$\frac{9}{20} \cdot \frac{6}{9} + \frac{11}{20} \cdot \frac{9}{11} = \frac{15}{20} = 0.75$$

No need to split if the leaf is pure
(all data have same labels)

Decision tree: accuracy

7

What is the train set accuracy now?

$$\frac{9}{20} \cdot \frac{6}{9} + \frac{11}{20} \cdot \frac{9}{11} = \frac{15}{20} = 0.75$$

Accuracy for two groups:

- Prereqs = yes (11): 9/11
- Prereqs = no (9): 6/9


For the 11 people prereqs = y, use the majority vote label **like** (9 like, 2 dislike).

Predicted label for 11 people is **like**, 9 people are correctly predicted.

			Prereqs	Lecturer	HasLabs		
			↓	↓	↓		
			Easy?	AI?	Sys?	Thy?	
Rating						Morning?	
consider it to be 'like'	+2	y	y	y	n	y	n
	+2	y	y	y	n	y	n
	+2	n	y	n	n	n	n
	+2	n	n	n	y	y	n
	+2	n	y	y	n	n	y
	+1	y	y	n	n	n	n
	+1	y	y	n	y	y	n
	+1	n	y	n	y	y	n
	0	n	n	n	n	n	y
	0	y	n	n	y	y	y
	0	n	y	n	y	y	n
	0	y	y	y	y	y	y
consider it to be 'dislike'	-1	y	y	y	n	y	y
	-1	n	n	y	y	y	n
	-1	n	n	y	n	n	y
	-1	y	n	y	n	n	y
	-2	n	n	y	y	y	n
	-2	n	y	y	n	n	y
	-2	y	n	y	n	n	n
	-2	y	n	y	n	n	y

KNN

- Select the number K of the neighbors
- Calculate the Euclidean distance of K number of neighbors
- Take the K nearest neighbors as per the calculated Euclidean distance.
- Among these k neighbors, count the number of the data points in each category.
- Assign the new data points to that category for which the number of the neighbor is maximum.



Training Data:

Person	height (feet)	weight (lbs)	foot size(inches)
male	6	180	12
male	5.92 (5'11")	190	11
male	5.58 (5'7")	170	12
male	5.92 (5'11")	165	10
female	5	100	6
female	5.5 (5'6")	150	8
female	5.42 (5'5")	130	7
female	5.75 (5'9")	150	9



Features

Task: Observe features x_1, \dots, x_D and predict class label $y \in \{1, \dots, C\}$

Naïve Bayes Model: Treat features as *conditionally independent* given class label,

$$p(x, y) = p(y)p(x|y) = p(y) \prod_{d=1}^D p(x_d | y)$$

build individual models for these

To classify a given instance x : Bayes rule!

$$p(y = c | x) = \frac{p(y = c)p(x | y = c)}{p(x)}$$

Key concept in Naïve Bayes

$$p(x, y) = p(y) p(x|y) = p(y) \prod_{d=1}^D p(x_d | y)$$

Class prior distribution

Class conditional distribution

Given one data point, it has 4 features (input), and the label is 0 (output)

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 0) &= p(y = 0) \cdot p(x_1, x_2, x_3, x_4 | y = 0) \\ &= p(y = 0) \cdot p(x_1 | y = 0) \cdot p(x_2 | y = 0) \cdot p(x_3 | y = 0) \cdot p(x_4 | y = 0) \end{aligned}$$

j: feature, c: label, i: data

$$y \sim \text{Categorical}(\pi_c) : p(y = c) = \pi_c$$

$$p(y = 1) = \pi_1$$

$$p(y = 2) = \pi_2$$

$$p(y = 3) = \pi_3 = 1 - \pi_1 - \pi_2$$

$$x|y \sim \text{Bernoulli}(\theta_{jc}) : p(x|y) = \theta_{jc}^x (1 - \theta_{jc})^{1-x}$$

$$x_{j=1}|y = 1 \sim \text{Bernoulli}(\theta_{j=1,c=1})$$

$$x_{j=2}|y = 1 \sim \text{Bernoulli}(\theta_{j=2,c=1})$$

$$x_{j=1}|y = 2 \sim \text{Bernoulli}(\theta_{j=1,c=2})$$

$$x_{j=2}|y = 2 \sim \text{Bernoulli}(\theta_{j=2,c=2})$$

$$x_{j=1}|y = 3 \sim \text{Bernoulli}(\theta_{j=1,c=3})$$

$$x_{j=2}|y = 3 \sim \text{Bernoulli}(\theta_{j=2,c=3})$$

y	x_1	x_2
1	0	1
3	1	0
3	1	1
2	0	0
1	1	0

Q: how many parameters?

$c-1+cj$

Model Selection and Evaluation

K-fold cross validation

- Randomly partition train set \mathcal{S} into K disjoint sets; call them $\text{fold}_1, \dots, \text{fold}_K$
- For each hyperparameter $h \in \{1, \dots, H\}$
 - For each $k \in \{1, \dots, K\}$
 - train \hat{f}_k^h with $\mathcal{S} \setminus \text{fold}_k$
 - measure error rate $e_{h,k}$ of \hat{f}_k^h on fold_k
 - Compute the average error of the above: $\widehat{err}^h = \frac{1}{K} \sum_{k=1}^K e_{h,k}$
- Choose $\hat{h} = \arg \min_h \widehat{err}^h$
- Train \hat{f}^* using \mathcal{S} (all the training points) with hyperparameter h
- Finally, evaluate \hat{f}^* on test set to estimate its future performance.

Use when (1) the dataset is small (2) ML algorithm's retraining time complexity is low (e.g., kNN)

5-fold cross validation



Q: If we use 5-fold cross validation for KNN, how many KNNs do we need to train?

A: 5 (if excluding last retraining step).

$$Error = \frac{1}{5} \sum_{i=1}^5 Error_i$$

		PREDICTED	
		POSITIVE	NEGATIVE
ACTUAL	POSITIVE	TRUE POSITIVES	FALSE NEGATIVES
	NEGATIVE	FALSE POSITIVES	TRUE NEGATIVES

Precision: dividing the true positives by anything that was predicted as a positive.

$$\frac{\text{TRUE POSITIVES}}{\text{TRUE POSITIVES} + \text{FALSE POSITIVES}}$$

		PREDICTED	
		POSITIVE	NEGATIVE
ACTUAL	POSITIVES	TRUE POSITIVES	FALSE NEGATIVES
	NEGATIVE	FALSE POSITIVES	TRUE NEGATIVES

Recall (or True Positive Rate): dividing the true positives by anything that should have been predicted as positive.

$$\frac{\text{TRUE POSITIVES}}{\text{TRUE POSITIVES} + \text{FALSE NEGATIVES}}$$

F1 score symmetrically represents both precision and recall in one metric.

$$F_1 = \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}} = \frac{\text{tp}}{\text{tp} + \frac{1}{2}(\text{fp} + \text{fn})}$$

- This is the *harmonic mean* of precision and recall
 - `harmonic_mean(x,y)`

$$\frac{1}{\frac{1}{2}(\frac{1}{x} + \frac{1}{y})}$$

- Gives equal importance to precision and recall – F1 may not be best when you care about one more than the other (e.g., in medical tests we care about recall)

Linear Models


Regression Learn a function that predicts outputs from inputs,

$$y = f(x)$$


Outputs y are real-valued

Linear Regression As the name suggests, uses a *linear function*:

$$y = w^T x + b$$



$$w^T x := \sum_{d=1}^D w_d x_d$$



Intuition Find a line that is as *close as possible* to every training data point


The distance from each point to the line is the **residual**

$$y - w^T x$$

Training Output

Prediction

Let's find w that will minimize the residual!



Functional Find a line that minimizes the sum of squared residuals!

Given: $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$

Compute:

$$w^* = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2$$

Least squares regression

Least squares can also be written more compactly,


$$\|x\| := \sqrt{x \cdot x}.$$

$$\min_w \sum_{i=1}^N (y^{(i)} - w^T x^{(i)})^2 = \|\mathbf{y} - \mathbf{X}w\|^2$$

Some slightly more advanced linear algebra gives us a solution,

$$w = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Ordinary Least Squares (OLS) solution OLS solution has less residual



Ordinary least-squares (OLS) estimation (no regularizer),

$$w^{\text{OLS}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2$$

$$\text{L2 norm: } \|w\| = \sqrt{\sum_{d=1}^D w_d^2}$$

$$\text{L1 norm: } \|w\|_1 = \sum_{d=1}^D |w_d|$$

L2-regularized Least-Squares (Ridge)

$$w^{\text{L2}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 + \lambda \|w\|^2$$

Convention: Just saying
'RLS' means L2-RLS

L1-regularized Least-Squares (LASSO)

LASSO: Least Absolute Shrinkage and Selection Operator

$$w^{\text{L2}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 + \lambda \|w\|_1$$

Model:


$$y \sim \text{Bernoulli}(p = \sigma(w^\top x))$$

Train: compute the MLE \hat{w}

Test: Given test point x^* compute

$$y^* = \arg \max_{v \in \{-1, 1\}} p(y = v \mid x^*; \hat{w})$$

- Equivalent to $y^* = \mathbf{I}\{\hat{w}^\top x^* \geq 0\}$



Nonlinear Models

A hyperplane $h(\mathbf{x})$ splits the original d -dimensional space into two half-spaces.
If the input dataset is linearly separable:


$$y = \begin{cases} +1 & \text{if } h(\mathbf{x}) > 0 \\ -1 & \text{if } h(\mathbf{x}) < 0 \end{cases}$$

Example:

$$h(x) = x_1 + 2x_2 - 4$$

Q: label for (0, 3)?

A: +1




Over all the n points, the **margin** of the linear classifier is the minimum distance of a point from the separating hyperplane:

$$\delta^* = \min_{\mathbf{x}_i} \left\{ \frac{y_i(\mathbf{w}^T \mathbf{x}_i + b)}{\|\mathbf{w}\|} \right\}$$

All the points that achieve this minimum distance are called **support vectors**.

$$\delta^* = \frac{y^*(\mathbf{w}^T \mathbf{x}^* + b)}{\|\mathbf{w}\|}$$



For training data $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$, a classifier $f(x) = w^\top x + b$ with 0 train error will satisfy

$$y^{(i)} f(x^{(i)}) = y^{(i)} (w^\top x^{(i)} + b) > 0$$

↓ negative margin when misclassifying it!

The distance for $(x^{(i)}, y^{(i)})$ to separating hyperplane

$$\frac{y^{(i)} (w^\top x^{(i)} + b)}{\|w\|}$$

The margin of a classifier $f(x)$ is

$$\min_i \frac{y^{(i)} (w^\top x^{(i)} + b)}{\|w\|}$$

Find f that maximize margin

$$\arg \max_{w, b} \min_i \frac{y^{(i)} (w^\top x^{(i)} + b)}{\|w\|}$$

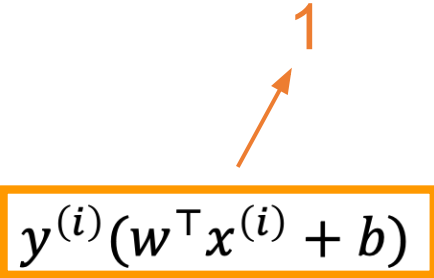
Way to solve this issue:

- Choose the scalar s such that the absolute distance of a **support vector** from the hyperplane is 1.

$$sy^*(\mathbf{w}^T \mathbf{x}^* + b) = 1$$


$$s = \frac{1}{y^*(\mathbf{w}^T \mathbf{x}^* + b)} = \frac{1}{y^* h(\mathbf{x}^*)}$$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \text{ for all points } \mathbf{x}_i \in \mathbf{D}$$

$$\arg \max_{\mathbf{w}, b} \min_i \frac{y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + b)}{\|\mathbf{w}\|}$$


Margin: $\delta^* = \frac{1}{\|\mathbf{w}\|}$

Max margin: $h^* = \arg \max_h \{\delta_h^*\} = \arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \right\}$




$$= \sigma \left(w_{1,0}a_0^{(0)} + w_{1,1}a_1^{(0)} + \dots + w_{1,n}a_n^{(0)} + b_1^{(0)} \right)$$

$$= \sigma \left(\sum_{i=1}^n w_{1,i}a_i^{(0)} + b_1^{(0)} \right)$$

$$\begin{pmatrix} a_1^{(1)} \\ a_2^{(1)} \\ \vdots \\ a_m^{(1)} \end{pmatrix} = \sigma \left[\begin{pmatrix} w_{1,0} & w_{1,1} & \dots & w_{1,n} \\ w_{2,0} & w_{2,1} & \dots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,0} & w_{m,1} & \dots & w_{m,n} \end{pmatrix} \begin{pmatrix} a_1^{(0)} \\ a_2^{(0)} \\ \vdots \\ a_n^{(0)} \end{pmatrix} + \begin{pmatrix} b_1^{(0)} \\ b_2^{(0)} \\ \vdots \\ b_m^{(0)} \end{pmatrix} \right]$$

$$a^{(1)} = \sigma \left(\mathbf{W}^{(0)} a^{(0)} + \mathbf{b}^{(0)} \right)$$

Number of parameters in this example: $m \times n + m$



$$784 \times 16 + 16 \times 16 + 16 \times 10$$

weights

$$16 + 16 + 10$$

biases

13,002

Each parameter has some impact
on the output...need to train all
these parameters simultaneously
to have a good prediction
accuracy

Clustering


Input: k : num. of clusters, $S = \{x_1, \dots, x_n\}$

[Initialize] Pick c_1, \dots, c_k as randomly selected points from S (see next slides for alternatives)

For $t=1, 2, \dots, \text{max_iter}$

- **[Assignments]** $\forall x \in S, \quad a_t(x) = \arg \min_{j \in [k]} \|x - c_j\|_2^2$
- If $t \neq 1$ AND $a_t(x) = a_{t-1}(x), \forall x \in S$
 - break
- **[Centroids]** $\forall j \in [k], \quad c_j \leftarrow \text{average}(\{x \in S: a_t(x) = j\})$

Output: c_1, \dots, c_k and $\{a_t(x_i)\}_{i \in [n]}$



Thank you!